

Database Technologies (TBD)

2021/2022

Database Query Optimization

Teaching Service

Henrique Reis Sendim Rodrigues (UP201606462)

João Carlos Machado Rocha Pires (UP201806079)

Mariana Catarina Pereira Soares (UP201605775)



FEUP FACULDADE DE ENGENHARIA
UNIVERSIDADE DO PORTO

Contents

1	Problem description	1
2	Setup	2
3	Questions	3
3.1	Question 1	3
3.1.1	SQL query	3
3.1.2	Answer	4
3.1.3	Execution Plans	4
3.1.4	Execution Time	5
3.2	Question 2	6
3.2.1	SQL query	6
3.2.2	Answer	6
3.2.3	Execution Plans	6
3.2.4	Execution Time	7
3.3	Question 3	7
3.3.1	a. Use not in.	8
3.3.1.1	SQL query	8
3.3.1.2	Answer	8
3.3.1.3	Execution Plans	9
3.3.1.4	Execution Time	11
3.3.2	b. Use external join and is null.	11
3.3.2.1	SQL query	11
3.3.2.2	Answer	11
3.3.2.3	Execution Plans	12
3.3.2.4	Execution Time	13
3.4	Question 4	13
3.4.1	SQL query	14
3.4.2	Answer	14
3.4.3	Execution Plans	14
3.4.4	Execution Time	14
3.5	Question 5	14

3.5.0.1	SQL query	14
3.5.0.2	Answer	14
3.5.0.3	Execution Plans	15
3.5.0.4	Execution Time	15
3.5.1	a. With a B-tree index on the type and academic year columns of the ZTI- POSAULA table;	15
3.5.1.1	Execution Plans	15
3.5.1.2	Execution Time	17
3.5.2	b. With a bitmap index on the type and academic year columns of the ZTI- POSAULA table.	17
3.5.2.1	Execution Plans	17
3.5.2.2	Execution Time	18
3.6	Question 6	18
3.6.1	SQL query	19
3.6.2	Answer	20
3.6.3	Execution Plans	20
3.6.4	Execution Time	21

1. Problem description

Our problem consists on a database with information regarding the distribution of teaching service in a faculty.

There are courses (table XUCS), described by a code (codigo), a designation (designacao), an acronym (sigla_uc) and a program (curso). Courses have occurrences in several years. Each occurrence is recorded by a row in the table XOCORRENCIAS, with information on the course code (codigo), academic year (ano_letivo), period of classes (periodo, that may be A-annual, 1S- first semester, 1T- first trimester, etc.), number of enrolled students (inscritos), students with distributed assessment (com_frequencia), number of approved (aprovados), course goals (objetivos) and content (conteudo), and department in charge (departamento).

Each occurrence may have one or more class types (T-theoretic, P-practical, L-laboratory, TP-theoretic/practical, OT- tutorial guidance). Each class type for an occurrence is recorded on table XTIPOSAULA with the number of similar classes (turnos), the number of week hours for each class (horas_turno), and in some cases the number of weekly classes (n_aulas).

The table XDSD records the teaching service distribution, in each semester, for each professor. More specifically, it records, for each class type of an occurrence, how many weekly hours are assigned to that professor. If a professor is teaching, in a single class, more than one course at the same time, for example from different programs, the weight of that course, in the perspective of the professor, may be less than 1 and recorded in attribute fator. Otherwise, the attribute fator will be 1. From the program perspective, the attribute fator is ignored. The attribute ordem enables listing the set of professors of a specific course occurrence in a specific order.

The professors are recorded in the table XDOCENTES with a number (nr), a name (nome), an acronym (sigla), a category code (categoria), a given name (proprio), a family name (apelido), and a status (estado: A-ativo, NA-não ativo, R-reformado).




Figure 1.1: Relational model for the case Teaching Service.

2. Setup

In the setup phase, we've copied the tables XDOCENTES, XDSD, XOCORRENCIAS, XTIPOSAULA and XUCS three times, each one with the prefixes 'X', 'Y' and 'Z'.

- X – no indexes and no integrity constraints.
- Y – with the standard integrity constraints (primary keys and foreign keys).
- Z – with the standard integrity constraints and extra indexes (the extra indexes were created as needed, which means they'll be justified in each question and their use is cumulative, starting from question 1 and ending up in question 6).

The following SQL code represents the copy process, followed by the creation of the standard integrity constraints for both 'Y' and 'Z' prefixed tables.

```
1 /* COPY TABLES - X, Y, Z */
2
3 CREATE TABLE XDOCENTES AS SELECT * FROM GTD10.xdocentes;
4 CREATE TABLE YDOCENTES AS SELECT * FROM GTD10.xdocentes;
5 CREATE TABLE ZDOCENTES AS SELECT * FROM GTD10.xdocentes;
6
7 CREATE TABLE XDSD AS SELECT * FROM GTD10.xdspd;
8 CREATE TABLE YDSD AS SELECT * FROM GTD10.xdspd;
9 CREATE TABLE ZDSD AS SELECT * FROM GTD10.xdspd;
10
11 CREATE TABLE XOCORRENCIAS AS SELECT * FROM GTD10.xocorrencias;
12 CREATE TABLE YOCORRENCIAS AS SELECT * FROM GTD10.xocorrencias;
13 CREATE TABLE ZOCORRENCIAS AS SELECT * FROM GTD10.xocorrencias;
14
15 CREATE TABLE XTIPOSAULA AS SELECT * FROM GTD10.xtiposaula;
16 CREATE TABLE YTIPOSAULA AS SELECT * FROM GTD10.xtiposaula;
17 CREATE TABLE ZTIPOSAULA AS SELECT * FROM GTD10.xtiposaula;
18
19 CREATE TABLE XUCS AS SELECT * FROM GTD10.xucs;
20 CREATE TABLE YUCS AS SELECT * FROM GTD10.xucs;
21 CREATE TABLE ZUCS AS SELECT * FROM GTD10.xucs;
22
23 /* CREATE PRIMARY AND FOREIGN KEY CONSTRAINTS Y */
24
25 ALTER TABLE YDOCENTES ADD CONSTRAINT YDOCENTES_PK PRIMARY KEY (NR);
26 ALTER TABLE YDSD ADD CONSTRAINT YDSD_PK PRIMARY KEY (NR, ID);
27 ALTER TABLE YOCORRENCIAS ADD CONSTRAINT YOCORRENCIAS_PK PRIMARY KEY (CODIGO, ANO_LETIVO,
    PERIODO);
28 ALTER TABLE YTIPOSAULA ADD CONSTRAINT YTIPOSAULA_PK PRIMARY KEY (ID);
29 ALTER TABLE YUCS ADD CONSTRAINT YUCS_PK PRIMARY KEY (CODIGO);
30
31 ALTER TABLE YDSD ADD CONSTRAINT YDSD_FK1 FOREIGN KEY (NR) REFERENCES YDOCENTES(NR);
```

```

32 ALTER TABLE YDSD ADD CONSTRAINT YDSD_FK2 FOREIGN KEY (ID) REFERENCES YTIPOSAULA(ID);
33 ALTER TABLE YOCORRENCIAS ADD CONSTRAINT YOCORRENCIAS_FK FOREIGN KEY (CODIGO) REFERENCES
    YUCS(CODIGO);
34 ALTER TABLE YTIPOSAULA ADD CONSTRAINT YTIPOSAULA_FK FOREIGN KEY (ANO_LETIVO, PERIODO,
    CODIGO) REFERENCES YOCORRENCIAS(ANO_LETIVO, PERIODO, CODIGO);
35
36 /* CREATE PRIMARY AND FOREIGN KEY CONSTRAINTS Z */
37
38 ALTER TABLE ZDOCENTES ADD CONSTRAINT ZDOCENTES_PK PRIMARY KEY (NR);
39 ALTER TABLE ZDSD ADD CONSTRAINT ZDSD_PK PRIMARY KEY (NR, ID);
40 ALTER TABLE ZOCORRENCIAS ADD CONSTRAINT ZOCORRENCIAS_PK PRIMARY KEY (CODIGO, ANO_LETIVO,
    PERIODO);
41 ALTER TABLE ZTIPOSAULA ADD CONSTRAINT ZTIPOSAULA_PK PRIMARY KEY (ID);
42 ALTER TABLE ZUCS ADD CONSTRAINT ZUCS_PK PRIMARY KEY (CODIGO);
43
44 ALTER TABLE ZDSD ADD CONSTRAINT ZDSD_FK1 FOREIGN KEY (NR) REFERENCES ZDOCENTES(NR);
45 ALTER TABLE ZDSD ADD CONSTRAINT ZDSD_FK2 FOREIGN KEY (ID) REFERENCES ZTIPOSAULA(ID);
46 ALTER TABLE ZOCORRENCIAS ADD CONSTRAINT ZOCORRENCIAS_FK FOREIGN KEY (CODIGO) REFERENCES
    ZUCS(CODIGO);
47 ALTER TABLE ZTIPOSAULA ADD CONSTRAINT ZTIPOSAULA_FK FOREIGN KEY (ANO_LETIVO, PERIODO,
    CODIGO) REFERENCES ZOCORRENCIAS(ANO_LETIVO, PERIODO, CODIGO);

```

3. Questions

In this chapter, each subsection will present the SQL query and the corresponding answer, as well as the execution plans and time for each one of the six raised questions. Each subsection starts with the question itself.

Note 1: It will only be presented the SQL query for the 'X' prefixed tables execution plan. The queries for the 'Y' and 'Z' prefixed tables execution plans are similar, with the difference of the prefix. The same happens to the result, which is the same for all, as expected.

Note 2: The execution times were calculated using the average of the obtained times for three consecutive executions on the same machine.

3.1 Question 1

Selection and join. Show the `codigo`, `designacao`, `ano_letivo`, `inscritos`, `tipo`, and `turnos` for the course 'Bases de Dados' of the program 275.

3.1.1 SQL query

The following SQL query selects the 'codigo', 'designacao', 'ano_letivo', 'inscritos', 'tipo' and 'turnos' fields from the join between the tables 'xucs', 'xocorrencias' and 'xtiposaula', with the constraint of the 'designacao' being equal to 'Bases de Dados' and 'curso' equal to 275.

```

1 SELECT xucs.codigo, xucs.designacao, xocorrencias.ano_letivo, xocorrencias.inscritos,
    xtiposaula.tipo, xtiposaula.turnos

```

```

2 FROM xucs JOIN xocorrencias ON xucs.codigo=xocorrencias.codigo JOIN xtiposaula ON
   xocorrencias.ano_letivo=xtiposaula.ano_letivo and xocorrencias.codigo=xtiposaula.
   codigo and xocorrencias.periodo=xtiposaula.periodo
3 WHERE xucs.designacao = 'Bases de Dados' and
4       xucs.curso = 275;

```

3.1.2 Answer

The result for this first question is the following:

	CODIGO	DESIGNACAO	ANO_LETIVO	INSCRITOS	TIPO	TURNOS
1	EIC3106	Bases de Dados	2003/2004	92	T	1
2	EIC3106	Bases de Dados	2003/2004	92	TP	4
3	EIC3106	Bases de Dados	2004/2005	114	T	1
4	EIC3106	Bases de Dados	2004/2005	114	TP	4
5	EIC3111	Bases de Dados	2005/2006	(null)	T	1
6	EIC3111	Bases de Dados	2005/2006	(null)	TP	6

Figure 3.1: Question 1 - Answer.

3.1.3 Execution Plans

Starting with the tables prefixed with 'X', the following picture shows the execution plan.

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				642
HASH JOIN				642
Access Predicates				
AND				
XOCORRENCIAS.ANO_LETIVO=XTIPOSADULA.ANO_LETIVO				
XOCORRENCIAS.CODIGO=XTIPOSADULA.CODIGO				
XOCORRENCIAS.PERIODO=XTIPOSADULA.PERIODO				
XUCS.CODIGO=XOCORRENCIAS.CODIGO				
MERGE JOIN		CARTESIAN	302	49
TABLE ACCESS	XUCS	FULL	1	13
Filter Predicates				
AND				
XUCS.DESIGNACAO='Bases de Dados'				
XUCS.CURSO=275				
BUFFER		SORT	21019	36
TABLE ACCESS	XTIPOSADULA	FULL	21019	36
TABLE ACCESS	XOCORRENCIAS	FULL	21747	593

Figure 3.2: Question 1 - X Execution Plan.

As we can see, the operation performed was to join the tables was a Semi Merge Join with the operation total cost being 642.

Moving on to the tables prefixed with 'Y', the following picture shows the execution plan, this time with a Nested Loops join and a total cost of 55.

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				55
HASH JOIN				55
Access Predicates				1
AND				
YOCORRENCIAS.ANO_LETIVO=YTIPOSAULA.ANO_LETIVO				
YOCORRENCIAS.CODIGO=YTIPOSAULA.CODIGO				
YOCORRENCIAS.PERIODO=YTIPOSAULA.PERIODO				
NESTED LOOPS			1	19
NESTED LOOPS			5	19
TABLE ACCESS	YUCS	FULL		13
Filter Predicates				1
AND				
YUCS.DESIGNACAO='Bases de Dados'				
YUCS.CURSO=275				
INDEX	YOCORRENCIAS_PK	RANGE SCAN	5	1
Access Predicates				
YUCS.CODIGO=YOCORRENCIAS.CODIGO				
TABLE ACCESS	YOCORRENCIAS	BY INDEX ROWID	5	6
TABLE ACCESS	YTIPOSAULA	FULL	21019	36

Figure 3.3: Question 1 - Y Execution Plan.

Last but not least, the execution plan for the tables prefixed with 'Z', similar to the one above, but with a lower cost since the following index was introduced, avoiding the 'and' operator between the two constraints:

```
1 CREATE INDEX ZUCS_DESIGNACAO ON ZUCS(DSIGNACAO);
```

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				45
HASH JOIN				45
Access Predicates				1
AND				
ZOCORRENCIAS.ANO_LETIVO=ZTIPOSAULA.ANO_LETIVO				
ZOCORRENCIAS.CODIGO=ZTIPOSAULA.CODIGO				
ZOCORRENCIAS.PERIODO=ZTIPOSAULA.PERIODO				
NESTED LOOPS			1	9
NESTED LOOPS			5	9
TABLE ACCESS	ZUCS	BY INDEX ROWID BATCHED		3
Filter Predicates				
ZUCS.CURSO=275				
INDEX	ZUCS_DESIGNACAO	RANGE SCAN	2	1
Access Predicates				
ZUCS.DESIGNACAO='Bases de Dados'				
INDEX	ZOCORRENCIAS_PK	RANGE SCAN	5	1
Access Predicates				
ZUCS.CODIGO=ZOCORRENCIAS.CODIGO				
TABLE ACCESS	ZOCORRENCIAS	BY INDEX ROWID	5	6
TABLE ACCESS	ZTIPOSAULA	FULL	21019	36

Figure 3.4: Question 1 - Z Execution Plan.

The reason behind the creation of this index was the fact that the 'designacao' field contains 3285 unique values, which might give us a performance advantage by having an index on it.

3.1.4 Execution Time

The execution times were 0,050 sec. for the first execution plan ('X') and 0,022 sec. for the last two ('Y' and 'Z').

3.2 Question 2

Aggregation. How many class hours of each type did the program 233 planned in year 2004/2005?

Note 3: For this question, we assumed that it was supposed to obtain the weekly hours for each type of class of the program 233, year 2004/2005, calculating by multiplying 'turnos' by 'horas_turno'.

3.2.1 SQL query

The following SQL query selects the 'tipo' and 'class_hours' (multiplication of 'turnos' by 'horas_turno') from the join of 'xucs' and 'xtiposaula' tables, having as constraints the 'ano_letivo' being equal to '2004/2005' and 'curso' equal to 233, with all this grouped by the type of class.

```
1 SELECT xtiposaula.tipo, sum(xtiposaula.turnos*xtiposaula.horas_turno) as "CLASS_HOURS"
2 FROM xucs JOIN xtiposaula ON xucs.codigo=xtiposaula.codigo
3 WHERE xtiposaula.ano_letivo = '2004/2005' and xucs.curso = 233
4 GROUP BY xtiposaula.tipo;
```

3.2.2 Answer

The following picture shows the answer for this question.

	TIPO	CLASS_HOURS
1	P	581,5
2	TP	697,5
3	T	308

Figure 3.5: Question 2 - Answer.

3.2.3 Execution Plans

The execution plans for this question use all Hash Joins. The difference is in the one for the 'Z' prefixed tables, since we created the following index, reducing the cost of the first two executions from 50 to 43.

```
1 CREATE INDEX ZUCS_CURSO ON ZUCS(CURSO);
```

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				50
HASH				50
HASH JOIN		GROUP BY		49
Access Predicates				54
XUCS.CODIGO=XTIPOSAULA.CODIGO				
TABLE ACCESS	XUCS	FULL		13
Filter Predicates				47
XUCS.CURSO=233				
TABLE ACCESS	XTIPOSAULA	FULL		36
Filter Predicates				1106
XTIPOSAULA.ANO_LETIVO='2004/2005'				

Figure 3.6: Question 2 - X Execution Plan.

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				50
HASH				50
HASH JOIN		GROUP BY		49
Access Predicates				53
ITEM_1=YTIPOSAULA.CODIGO				
VIEW	SYS.VW_GBF_5			13
TABLE ACCESS	YUCS	FULL		47
Filter Predicates				47
YUCS.CURSO=233				
TABLE ACCESS	YTIPOSAULA	FULL		36
Filter Predicates				1106
YTIPOSAULA.ANO_LETIVO='2004/2005'				

Figure 3.7: Question 2 - Y Execution Plan.

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				43
HASH				43
HASH JOIN		GROUP BY		42
Access Predicates				53
ITEM_1=ZTIPOSAULA.CODIGO				
VIEW	SYS.VW_GBF_5			6
TABLE ACCESS	ZUCS	BY INDEX ROWID BATCHED		47
INDEX	ZUCS_CURSO	RANGE SCAN		6
Access Predicates				1
ZUCS.CURSO=233				
TABLE ACCESS	ZTIPOSAULA	FULL		36
Filter Predicates				1106
ZTIPOSAULA.ANO_LETIVO='2004/2005'				

Figure 3.8: Question 2 - Z Execution Plan.

3.2.4 Execution Time

Regarding the execution time, it was the same for all the execution plans: 0,020 sec.

3.3 Question 3

Negation. Which courses (show the code) did have occurrences planned but did not get service assigned in year 2003/2004?

Note 4: This question raised some ambiguity inside the group, so we have decided to assume it was referring to something in particular: we assumed that it was intended to get the 'codigo' of the UCS that didn't had any 'inscritos' in the 'ano_letivo' equal to '2003/2004'. We also considered to

include UCS with 'inscritos', but without 'com_frequencia'. However, this last option was discarded since there was no way to check when we were in the presence of a non-occurrence or of an occurrence without a single student having enough frequency.

3.3.1 a. Use not in.

3.3.1.1 SQL query

For the first part of the question, this SQL query select the 'codigo' from 'xocorrencias' where the same 'codigo' is not in the list of 'codigo' obtained from a sub-query where the 'inscritos' field is not null and the 'ano_letivo' is '2003/2004'.

```
1 SELECT CODIGO
2 FROM XOCORRENCIAS
3 WHERE CODIGO NOT IN (SELECT CODIGO FROM XOCORRENCIAS WHERE INSCRITOS IS NOT NULL AND
    ANO_LETIVO = '2003/2004') AND ANO_LETIVO = '2003/2004';
```

3.3.1.2 Answer

The answer for this first part of the question was different of the answer obtained for the second part. Having searched for the reason why that happened, we found that the 'NOT IN' in Oracle SQL has a problem with 'NULL' values, which makes it produce different results.

	⚡ CODIGO		⚡ CODIGO
1	MEM170	31	EIC5125
2	EIC4221	32	MTM114
3	MEM157	33	EC5287
4	EIC4222	34	MEM1209
5	EIC4220	35	MEM180
6	MEM181	36	EIC5122
7	MEM189	37	EIC5123
8	EIC4225	38	MEM184
9	EMG5206	39	EIC4223
10	MEM187	40	MFAMF1204
11	MEM132	41	EI1208
12	EI1107	42	MEM163
13	EIC5127	43	MEM153
14	MEM164	44	MEM191
15	MEA317	45	MMI1203
16	MEM179	46	MEM173
17	MTM117	47	MFAMF1108
18	MTM111	48	MEM188
19	MEM1208	49	MEEC2105
20	MFAMF1106	50	MMI1104
21	MFAMF1207	51	MTM104
22	EEC5225	52	EIC4224
23	MEA212	53	MEM1117
24	EIC5124	54	EIC5129
25	MEAM5000	55	MMCCE1200
26	EMM528	56	MEM183
27	MEA402	57	MEM167
28	MEEC1089	58	MEM158
29	MEM111	59	EI1103
30	EIC5126	60	MEM182

Figure 3.9: Question 3a. - Answer.

3.3.1.3 Execution Plans

For this question, having the two indexes previously defined, we didn't found the need to create new indexes. Also, for all the execution plans, a Hash Join was verified. The cost for the first was 1185, 595 for the second and third.

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				152 1185
HASH JOIN		ANTI		152 1185
Access Predicates CODIGO=CODIGO				
TABLE ACCESS	YOCORRENCIAS	FULL		483 593
Filter Predicates ANO_LETIVO='2003/2004'				
TABLE ACCESS	YOCORRENCIAS	FULL		326 593
Filter Predicates AND ANO_LETIVO='2003/2004' INSCRITOS IS NOT NULL				

Figure 3.10: Question 3a. - X Execution Plan.

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				152 595
HASH JOIN		ANTI		152 595
Access Predicates CODIGO=CODIGO				
NESTED LOOPS		ANTI		152 595
STATISTICS COLLECTOR				
INDEX	YOCORRENCIAS_PK	FAST FULL SCAN		483 27
Filter Predicates ANO_LETIVO='2003/2004'				
TABLE ACCESS	YOCORRENCIAS	BY INDEX ROWID BATCHED		224 568
Filter Predicates INSCRITOS IS NOT NULL				
INDEX	YOCORRENCIAS_PK	RANGE SCAN		483 93
Access Predicates AND CODIGO=CODIGO ANO_LETIVO='2003/2004'				
TABLE ACCESS	YOCORRENCIAS	BY INDEX ROWID BATCHED		326 568
Filter Predicates INSCRITOS IS NOT NULL				
INDEX	YOCORRENCIAS_PK	SKIP SCAN		483 93
Access Predicates ANO_LETIVO='2003/2004'				
Filter Predicates ANO_LETIVO='2003/2004'				

Figure 3.11: Question 3a. - Y Execution Plan.

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				595
HASH JOIN		ANTI		595
Access Predicates CODIGO=CODIGO				152
NESTED LOOPS		ANTI		595
STATISTICS COLLECTOR				152
INDEX	ZOCORRENCIAS_PK	FAST FULL SCAN	483	27
Filter Predicates ANO_LETIVO='2003/2004'				
TABLE ACCESS	ZOCORRENCIAS	BY INDEX ROWID BATCHED	224	568
Filter Predicates INSCRITOS IS NOT NULL				
INDEX	ZOCORRENCIAS_PK	RANGE SCAN	483	93
Access Predicates AND CODIGO=CODIGO ANO_LETIVO='2003/2004'				
TABLE ACCESS	ZOCORRENCIAS	BY INDEX ROWID BATCHED	326	568
Filter Predicates INSCRITOS IS NOT NULL				
INDEX	ZOCORRENCIAS_PK	SKIP SCAN	483	93
Access Predicates ANO_LETIVO='2003/2004'				
Filter Predicates ANO_LETIVO='2003/2004'				

Figure 3.12: Question 3a. - Z Execution Plan.

3.3.1.4 Execution Time

Regarding the execution times, they were similar to all the execution plans: 0,046 sec.

3.3.2 b. Use external join and is null.

3.3.2.1 SQL query

For this second part of the question, instead of the use of 'NOT IN', the 'codigo' field is selected from the join between the tables 'xucs' and 'xocorrencias' considering all the entries where 'inscritos' is 'NULL' and for the 'ano_letivo' equal to '2003/2004'.

```

1 SELECT XUCS.CODIGO
2 FROM XUCS JOIN XOCORRENCIAS ON XUCS.CODIGO = XOCORRENCIAS.CODIGO
3 WHERE XOCORRENCIAS.INSCRITOS IS NULL AND XOCORRENCIAS.ANO_LETIVO = '2003/2004';

```

3.3.2.2 Answer

The answer, as mentioned above, was different.

	CODIGO		CODIGO
1	EC5287	34	MEM132
2	EEC5225	35	MEM153
3	EI1103	36	MEM157
4	EI1107	37	MEM158
5	EI1208	38	MEM163
6	EIC4220	39	MEM164
7	EIC4221	40	MEM167
8	EIC4222	41	MEM170
9	EIC4223	42	MEM173
10	EIC4224	43	MEM175
11	EIC4225	44	MEM179
12	EIC5122	45	MEM180
13	EIC5123	46	MEM181
14	EIC5124	47	MEM182
15	EIC5125	48	MEM183
16	EIC5126	49	MEM184
17	EIC5127	50	MEM187
18	EIC5129	51	MEM188
19	EM446	52	MEM189
20	EMG5206	53	MEM191
21	EMM528	54	MFAMF...
22	GEI2414	55	MFAMF...
23	MEA111	56	MFAMF...
24	MEA212	57	MFAMF...
25	MEA317	58	MGI1204
26	MEA402	59	MMCCE...
27	MEAM5000	60	MMI1104
28	MEEC1089	61	MMI1203
29	MEEC2105	62	MRSC1104
30	MEM111	63	MTM104
31	MEM1117	64	MTM111
32	MEM1208	65	MTM114
33	MEM1209	66	MTM117

Figure 3.13: Question 3a. - Answer.

3.3.2.3 Execution Plans

Regarding the execution plans, for the first case, a Hash Join was used (cost = 606), while for the remaining ones it was a Table Access by RowID (cost = 568).

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				157 606
HASH JOIN				157 606
Access Predicates XUCS.CODIGO=XOCORRENCIAS.CODIGO				
TABLE ACCESS	XOCORRENCIAS	FULL		157 593
Filter Predicates AND XOCORRENCIAS.INSCRITOS IS NULL XOCORRENCIAS.ANO_LETIVO='2003/2004'				
TABLE ACCESS	XUCS	FULL		5396 13

Figure 3.14: Question 3b. - X Execution Plan.

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				157 568
TABLE ACCESS	YOCORRENCIAS	BY INDEX ROWID BATCHED		157 568
Filter Predicates YOCORRENCIAS.INSCRITOS IS NULL				
INDEX	YOCORRENCIAS_PK	SKIP SCAN		483 93
Access Predicates YOCORRENCIAS.ANO_LETIVO='2003/2004'				
Filter Predicates YOCORRENCIAS.ANO_LETIVO='2003/2004'				

Figure 3.15: Question 3b. - Y Execution Plan.

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				157 568
TABLE ACCESS	ZOCORRENCIAS	BY INDEX ROWID BATCHED		157 568
Filter Predicates ZOCORRENCIAS.INSCRITOS IS NULL				
INDEX	ZOCORRENCIAS_PK	SKIP SCAN		483 93
Access Predicates ZOCORRENCIAS.ANO_LETIVO='2003/2004'				
Filter Predicates ZOCORRENCIAS.ANO_LETIVO='2003/2004'				

Figure 3.16: Question 3b. - Z Execution Plan.

3.3.2.4 Execution Time

Again, as for the first part of the question, the execution times were similar: 0,046 sec. ('X') and 0,042 sec. ('Y' and 'Z').

3.4 Question 4

Who is the professor with more class hours for each type of class, in the academic year 2003/2004? Show the number and name of the professor, the type of class and the total of class hours times the factor.

3.4.1 SQL query

3.4.2 Answer

3.4.3 Execution Plans

3.4.4 Execution Time

3.5 Question 5

Compare the execution plans (just the environment Z) and the index sizes for the query giving the course code, the academic year, the period, and number of hours of the type 'OT' in the academic years of 2002/2003 and 2003/2004.

Note 5: This question raised some ambiguity inside the group, so we have decided to assume it was referring to something in particular: we assumed that it was intended to compare the execution plan of the query without indexes with the query using each type of indexes. We also assumed that it was supposed to compare the index size between the indexes of the type and academic year columns.

3.5.0.1 SQL query

Note 6: For this question, we assumed that it was supposed to obtain the weekly hours for the type of class equal to 'OT' of the year 2002/2003 or 2003/2004, calculating by multiplying 'turnos' by 'horas_turno'.

The following SQL query selects the 'codigo', 'ano_letivo', 'periodo' and 'num_horas' (multiplication of 'turnos' by 'horas_turno') from the join of 'zucs' and 'ztiposaula' tables, having as constraints the 'tipo' equal to 'OT' and the 'ano_letivo' equal to 2002/2003 or 2003/2004, with all this grouped by course code, academic year and period.

```
1 SELECT ZUCS.CURSO, ZTIPOSAULA.ANO_LETIVO, ZTIPOSAULA.PERIODO, SUM(ZTIPOSAULA.TURNOS*
   ZTIPOSAULA.HORAS_TURNO) AS "NUM_HORAS"
2 FROM ZUCS JOIN ZTIPOSAULA ON ZUCS.CODIGO=ZTIPOSAULA.CODIGO
3 WHERE (ZTIPOSAULA.ANO_LETIVO = '2002/2003' OR ZTIPOSAULA.ANO_LETIVO = '2003/2004')
4 AND ZTIPOSAULA.TIPO = 'OT'
5 GROUP BY ZUCS.CURSO, ZTIPOSAULA.ANO_LETIVO, ZTIPOSAULA.PERIODO;
```

3.5.0.2 Answer

The following picture shows the answer for this question.

	CURSO	ANO_LETIVO	PERIODO	NUM_HORAS
1	275	2002/2003	2S	27
2	275	2003/2004	2S	24

Figure 3.17: Question 5. - Answer.

3.5.0.3 Execution Plans

The following picture shows the execution plan for the above query.

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				50
HASH				50
HASH JOIN		GROUP BY	443	49
Access Predicates				
ZUCS.CODIGO=ZTIPOSAULA.CODIGO				
NESTED LOOPS			443	36
NESTED LOOPS				
STATISTICS COLLECTOR				
TABLE ACCESS	ZTIPOSAULA	FULL	443	
Filter Predicates				
AND				
ZTIPOSAULA.TIPO='OT'				
OR				
ZTIPOSAULA.ANO_LETIVO='2002/2003'				
ZTIPOSAULA.ANO_LETIVO='2003/2004'				
INDEX	ZUCS_PK	UNIQUE SCAN		
Access Predicates				
ZUCS.CODIGO=ZTIPOSAULA.CODIGO				
TABLE ACCESS	ZUCS	BY INDEX ROWID	1	13
TABLE ACCESS	ZUCS	FULL	5396	13

Figure 3.18: Question 5. - Z Execution Plan.

As we can see, the operation total cost was 50.

3.5.0.4 Execution Time

The execution time was 0,027 sec.

3.5.1 a. With a B-tree index on the type and academic year columns of the ZTIPOSAULA table;

3.5.1.1 Execution Plans

Before running the query presented in question 5, we created the following B-tree indexes.

```
1 CREATE INDEX ZTIPOSAULA_TIPO ON ZTIPOSAULA(TIPO);
2 CREATE INDEX ZTIPOSAULA_ANOLETIVO ON ZTIPOSAULA(ANO_LETIVO);
```

Note 7: By default, Oracle Database creates B-tree indexes.

The following picture shows the execution plan for the query.

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				28
HASH				19
HASH JOIN		GROUP BY		28
Access Predicates				18
ZUCS.CODIGO=ZTIPOSAULA.CODIGO				
NESTED LOOPS				28
STATISTICS COLLECTOR				18
TABLE ACCESS	ZTIPOSAULA	BY INDEX ROWID BATCHED		5
Filter Predicates				
OR				
ZTIPOSAULA.ANO_LETIVO='2002/2003'				
ZTIPOSAULA.ANO_LETIVO='2003/2004'				
INDEX	ZTIPOSAULA_TIPO	RANGE SCAN		1
Access Predicates				
ZTIPOSAULA.TIPO='OT'				
INDEX	ZUCS_PK	UNIQUE SCAN		
Access Predicates				
ZUCS.CODIGO=ZTIPOSAULA.CODIGO				
TABLE ACCESS	ZUCS	BY INDEX ROWID		13
TABLE ACCESS	ZUCS	FULL		5396

Figure 3.19: Question 5a. - Z Execution Plan.

As we can see, by using a index in type column, the operation total cost decreased from 50 to 19.

Note 8: By default, Oracle Database chooses if it should use the indexes or not in order to minimize costs. In this case, it choose not to use the index 'ZTIPOSAULA_ANOLETIVO'. If we force to use both indexes, using the hint '/*+INDEX(ZTIPOSAULA ZTIPOSAULA_ANOLETIVO)*/' immediately after the select instruction, we verified that the cost goes up to 75.

To get the size of the index used in the column type we used the following query.

```
1 select sum(bytes)/1024/1024 as "Index Size (MB)" from user_segments where segment_name='
  ZTIPOSAULA_TIPO';
```

As a result of that query, we obtained the following size to the 'ZTIPOSAULA_TIPO' index.

	Index Size (MB)
1	0,375

Figure 3.20: Index Size Query - Answer.

To get the size of the index used in the column academic year we used the following query.

```
1 select sum(bytes)/1024/1024 as "Index Size (MB)" from user_segments where segment_name='
  ZTIPOSAULA_ANOLETIVO';
```

As a result of that query, we obtained the following size to the 'ZTIPOSAULA_ANOLETIVO' index.

	Index Size (MB)
1	0,5625

Figure 3.21: Index Size Query - Answer.

We can see that the index size created in the academic year column is greater than the one created in the type column. This could be the reason so that the 'ZTIPOSAULA_ANOLETIVO' index is not used by the oracle optimizer.

3.5.1.2 Execution Time

The execution time was 0,024 sec.

3.5.2 b. With a bitmap index on the type and academic year columns of the ZTIPOSAULA table.

3.5.2.1 Execution Plans

Before running the query presented in question 5, we created the following bitmap indexes.

```
1 CREATE BITMAP INDEX BM_ZTIPOSAULA_TIPO ON ZTIPOSAULA(TIPO);
2 CREATE BITMAP INDEX BM_ZTIPOSAULA_ANOLETIVO ON ZTIPOSAULA(ANO_LETIVO);
```

The following picture shows the execution plan for the query.

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				28
HASH		GROUP BY		22
HASH JOIN				21
Access Predicates ZUCS.CODIGO=ZTIPOSAULA.CODIGO				
NESTED LOOPS				28
NESTED LOOPS				21
STATISTICS COLLECTOR				
TABLE ACCESS	ZTIPOSAULA	BY INDEX ROWID BATCHED TO ROWIDS		8
BITMAP CONVERSION				
BITMAP AND				
BITMAP INDEX	BM_ZTIPOSAULA_TIPO	SINGLE VALUE		
Access Predicates ZTIPOSAULA.TIPO='OT'				
BITMAP OR				
BITMAP INDEX	BM_ZTIPOSAULA_ANOLETIVO	SINGLE VALUE		
Access Predicates ZTIPOSAULA.ANO_LETIVO='2002/2003'				
BITMAP INDEX	BM_ZTIPOSAULA_ANOLETIVO	SINGLE VALUE		
Access Predicates ZTIPOSAULA.ANO_LETIVO='2003/2004'				
INDEX	ZUCS_PK	UNIQUE SCAN		
Access Predicates ZUCS.CODIGO=ZTIPOSAULA.CODIGO				
TABLE ACCESS	ZUCS	BY INDEX ROWID	1	13
TABLE ACCESS	ZUCS	FULL	5396	13


Figure 3.22: Question 5b. - Z Execution Plan.

As we can see, by using bitmap indexes in the mentioned columns in the question, the total cost decreased from 50 to 22.

To get the index size of the type column we used the following query.

```
1 select sum(bytes)/1024/1024 as "Index Size (MB)" from user_segments
2 where segment_name='BM_ZTIPOSAULA_TIPO';
```

As a result of that query, we obtained the following size to the 'BM_ZTIPOSAULA_TIPO' index.



	Index Size (MB)
1	0,0625

Figure 3.23: Index Size Query - Answer.

To get the index size of the academic year column we used the following query.

```
1 select sum(bytes)/1024/1024 as "Index Size (MB)" from user_segments where segment_name='
BM_ZTIPOSAULA_ANOLETIVO';
```

As a result of that query, we obtained the following size to the 'BM_ZTIPOSAULA_ANOLETIVO' index.



	Index Size (MB)
1	0,0625

Figure 3.24: Index Size Query - Answer.

The size of each index turned out to be the same and smaller than in the previous ones.

3.5.2.2 Execution Time

The execution time was 0,024 sec.

3.6 Question 6

Select the programs (curso) that have classes with all the existing types.

Note 9: For this question, we found ourselves trying to understand which of the following 3 interpretations was the correct one for this question:

- A. Get the UCs with classes from all types.

- B. Get the courses with classes from all types, i.e. the UCs belonging to a course may not have classes from all types, but all the classes, for all the UCs of that course, would have (or not) the overall class type representatives.
- C. Get the name(s) of the course(s) with at least 1 UC with all the class types.

We've decided that the interpretation B. was the one desired, so we moved one with that one in what respects to the analysis of the answer and the execution plan and time. However, before proceeding with that option, the following two SQL queries represent the ones that would retrieve the results for the interpretations A. and C.

```

1 /*A*/
2 select codigo
3 from (select codigo, tipo
4       from xtiposaula
5       group by codigo, tipo)
6 group by codigo
7 having count(codigo) = (select count(*)
8                        from (select distinct tipo
9                             from xtiposaula));

1 /*C*/
2 create view aux_codigo as
3 select codigo
4 from (select codigo, tipo
5       from xtiposaula
6       group by codigo, tipo)
7 group by codigo
8 having count(codigo) = (select count(*)
9                        from (select distinct tipo
10                             from xtiposaula));
11
12 select distinct curso
13 from aux_codigo join xucs on xucs.codigo = aux_codigo.codigo;
```

3.6.1 SQL query

The following SQL query, for the interpretation B., selects the 'curso' field from a sub-query that gets the list of courses and the corresponding class types, grouping then by course and having as constraint the fact that the course carnality (in this case the number of different class types) should be the same as the number of different course types (5) obtained by a nested sub-query.

```

1 /*B*/
2 select curso
3 from (select curso, tipo
4       from xtiposaula join xucs on xtiposaula.codigo = xucs.codigo
5       group by curso, tipo)
6 group by curso
7 having count(curso) = (select count(*)
8                        from (select distinct tipo
9                             from xtiposaula));
```

3.6.2 Answer

The answer for this question is the one that follows:

	CURSO
1	9461
2	4495
3	9508
4	2021

Figure 3.25: Question 6. - Answer.

3.6.3 Execution Plans

Considering the indexes created until now, we didn't find the need to create a new one for this question in particular.

That said, the following three pictures represent the three different execution plans for the 'X', 'Y' and 'Z' prefixed tables.

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				51
FILTER			2	
Filter Predicates				
COUNT(\$vm_col_1) = (SELECT COUNT(*) FROM (SELECT DISTINCT TIPO TIPO FROM XTIPOSAULA XTIPOSAULA) from\$query_005)				
HASH		GROUP BY	2	51
VIEW	SYS.VM_NWWW_0		404	51
HASH		GROUP BY	404	51
HASH JOIN			21019	49
Access Predicates				
XTIPOSAULA.CODIGO=XUCS.CODIGO				
TABLE ACCESS	XUCS	FULL	5396	13
TABLE ACCESS	XTIPOSAULA	FULL	21019	36
SORT		AGGREGATE	1	
VIEW			5	37
SORT		UNIQUE	5	37
TABLE ACCESS	XTIPOSAULA	FULL	21019	36

Figure 3.26: Question 6 - X Execution Plan.

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				2 51
FILTER				
Filter Predicates	COUNT(\$vm_col_1) = (SELECT COUNT(*) FROM (SELECT DISTINCT TIPO TIPO FROM YTIPOSAULA YTIPOSAULA) from\$_subquery\$_005)			
HASH		GROUP BY	2	51
VIEW	SYS.VM_NWWW_Q		404	51
HASH		GROUP BY	404	51
HASH JOIN			21019	49
Access Predicates	YTIPOSAULA.CODIGO=YUCS.CODIGO			
TABLE ACCESS	YUCS	FULL	5396	13
TABLE ACCESS	YTIPOSAULA	FULL	21019	36
SORT		AGGREGATE	1	
VIEW			5	37
SORT		UNIQUE	5	37
TABLE ACCESS	YTIPOSAULA	FULL	21019	36

Figure 3.27: Question 6 - Y Execution Plan.

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				2 51
FILTER				
Filter Predicates	COUNT(\$vm_col_1) = (SELECT COUNT(*) FROM (SELECT DISTINCT TIPO TIPO FROM ZTIPOSAULA ZTIPOSAULA) from\$_subquery\$_005)			
HASH		GROUP BY	2	51
VIEW	SYS.VM_NWWW_Q		404	51
HASH		GROUP BY	404	51
HASH JOIN			21019	49
Access Predicates	ZTIPOSAULA.CODIGO=ZUCS.CODIGO			
TABLE ACCESS	ZUCS	FULL	5396	13
TABLE ACCESS	ZTIPOSAULA	FULL	21019	36
SORT		AGGREGATE	1	
VIEW			5	37
SORT		UNIQUE	5	37
TABLE ACCESS	ZTIPOSAULA	FULL	21019	36

Figure 3.28: Question 6 - Z Execution Plan.

As we can see, the Hash Join was used for the three cases, and the cost was also the same for all: 51.

3.6.4 Execution Time

Regarding the execution times, they were 0,032, 0,030 and 0,029 sec. for the 'X', 'Y' and 'Z' prefixed tables, respectively.